Interim UTKFace test results: 64/66 categorical main runs complete
50 full-data epochs; mean ± sample SD across available seeds; checkpoints selected by validation MAE

Stat MaxRL MAE, width 2 (population) (n=3)
TailRL (64 samples) (n=1)
TailRL (population) (n=3)
Direct MAE (categorical mean) (n=3)
Stat MaxRL MAE, width 2 (64 samples) (n=3)
Stat MaxRL MAE, width 5 (64 samples) (n=3)
Direct MSE (categorical mean) (n=3)
Stat MaxRL MAE, width 5 (population) (n=3)
Stat MaxRL MSE, width 2 (population) (n=3)
Stat MaxRL MSE, width 2 (64 samples) (n=3)
Binary MaxRL (population) (n=3)
Binary MaxRL (64 samples) (n=3)
Stat MaxRL MSE, width 5 (population) (n=3)
Stat MaxRL MSE, width 5 (64 samples) (n=3)
GRPO MAE (64 samples) (n=3)
GRPO MAE (population) (n=3)
RLOO MAE (64 samples) (n=3)
RLOO MAE (population) (n=3)
GRPO MSE (64 samples) (n=3)
RLOO MSE (population) (n=3)
RLOO MSE (64 samples) (n=3)
GRPO MSE (population) (n=3)

4.5 5.0 5.5 6.0 6.5 7.0
Test mean-prediction MAE (years; lower is better)

0.880 0.885 0.890 0.895 0.900 0.905 0.910
Test median Spearman (higher is better)

0 10 20 30 40 50 60
Test label NLL (nats; lower is better)

Multiple seeds; bars show sample SD
One seed; uncertainty unavailable